

Nifty 100 Financial Intelligence Platform

Acceptance Checklist -- 20 Quality Gates

This is an honest verification record, not a sign-off formality. Each criterion below was checked against real data/files/running services at the time of writing -- not assumed correct because the underlying code exists or ran without error. PARTIAL and FAIL verdicts are reported as found; none were adjusted for appearance. Summary: 15 PASS (2 with a caveat, marked PASS*), 3 PARTIAL, 1 FAIL, 1 UNMEASURED.

ID	Criterion	Verdict	Notes
AC-01	92 companies present, no extra/missing tickers	PASS	SELECT COUNT(*) FROM companies = 92, verified directly.
AC-02	≥ 90% of companies have ≥ 10yr P&L/BS/CF records	PASS	93.4-95.7% across all three tables.
AC-03	All FK relationships intact (PRAGMA foreign_key_check = 0 rows)	PASS	0 violations, verified directly.
AC-04	financial_ratios ≥ 1,100 rows, all 14 KPI columns populated	FAIL	1,070 rows (short by 30) and 13 declared columns, not 14 -- financial_ratios.xlsx's own schema (spec Section 6.4) never declared a 14th column. Each column is 97.5-99.9% populated, not fully null-free.
AC-05	Revenue CAGR for 3 companies matches hand-computed value ± 0.1%	PASS	Recomputed independently from raw sales figures via SQL, not by calling the engine's own code.
AC-06	ROE for 5 companies matches companies.roe_percentage ± 5%	PASS*	4/5 within 3 percentage points. TCS is off by 50pp, but the spec itself (Section 5.1) flags companies.roe_percentage=0.52 for TCS as a known anomaly and says to trust the Ratio Engine value instead -- which is what this platform does.
AC-07	Quality preset screener produces 10-50 companies	PASS	38 companies matched.
AC-08	Company Profile screen loads in < 3 seconds	UNMEASURED	Qualitatively fast in every manual test across every sprint; no automated click-to-render timing harness available in this environment to produce a rigorous number.
AC-09	Screener CSV download is well-formed with correct headers	PASS	Round-trips cleanly through a full write/parse cycle; headers match exactly.
AC-10	No text overflow or overlapping pages in generated reports	PARTIAL	Text-extraction check on 5 random tearsheets confirms correct 2-page structure and real content on every page. Cannot visually confirm absence of overlap/overflow -- no PDF renderer (poppler) available in this environment.
AC-11	GET /api/v1/health returns HTTP 200 with db_row_counts for all 10 tables	PARTIAL	Returns 200 with row counts for all 12 real tables. The database has 12 tables, not 10 -- same discrepancy documented in db/schema.sql since Sprint 1 (the spec's own "10 tables" claim doesn't match its own entity-relationship map, which names 12).
AC-12	GET /api/v1/companies/TCS/ratios returns rows for ≥ 10 years	PASS	12 rows returned.
AC-13	GET /api/v1/screener?min_roe=15&max_de=1 is consistent with Module 3 output	PASS	API and the underlying screener engine agree exactly (49 companies). No sheet in screener_output.xlsx has this exact 2-filter combination pre-computed to diff against directly (the Quality preset adds a 3rd filter, FCF), but the live consistency check confirms the same logic drives both.
AC-14	Peer percentile table populated for all 11 peer groups, all members present	PASS	11/11 groups; all 56 members present in both peer_groups and peer_percentiles.

ID	Criterion	Verdict	Notes
AC-15	All 92 companies assigned to a cluster (0-4), no nulls	PASS	Verified directly: 0 nulls, all 92 assigned.
AC-16	pros_cons_generated.csv has ≥ 1 pro and ≥ 1 con for every company	PASS	0 companies missing either a pro or a con.
AC-17	92 tearsheet PDFs exist, each ≥ 50 KB	PASS	All 92 present, all above the size threshold.
AC-18	pytest shows ≥ 60 tests collected, 0 failures	PASS	169 tests passing, well over the threshold.
AC-19	validation_failures.csv: every row has company_id, field, issue, severity	PARTIAL	1 of 317 rows (a DQ-16 coverage warning for JIOFIN) has a null field value -- DQ-16 measures year-count across a company's whole history, which doesn't naturally map to one column.
AC-20	analyst_guide.pdf exists, ≥ 10 pages, covers screener and dashboard	PASS	Exactly 10 pages. Both required sections present, plus peer comparison, valuation, API reference, KPI glossary, and a dedicated data-quality-caveats section.

Notes on the 3 Deliverables-Tracker Path Mismatches

During the D-01 through D-23 deliverables verification pass (separate from this 20-criteria check), 3 files were found to exist with real content but at a different path than the deliverables tracker states: D-04 (exploratory_queries.sql is at db/, not notebooks/), D-17 (only 10 sector report PDFs exist, matching the 10 real broad sectors in sectors.xlsx -- the tracker's "11" does not match the real data, same gap documented since Sprint 1's sector distribution finding), and D-22 (analyst_guide.pdf is at reports/, not docs/). None of these are missing deliverables -- all exist with substantive, verified content -- but the tracker's stated paths/counts don't match where the work actually landed.

Also found during that pass: the "simulated data must be labeled SIMULATED" rule is only partially satisfied -- 2 of 4 dashboard screens that display market_cap-derived data (Sector Analysis, Valuation) show the notice; 2 (Home, Screener) do not.